

Embedded Systems and IoT

Ingegneria e Scienze Informatiche - UNIBO

a.a 2022/2023

Lecturer: Prof. Alessandro Ricci

[module-2.3]

TASK-BASED ARCHITECTURES

OUTLINE

- Task-based Architecture
 - integration with synchronous FSM
- Cooperative scheduling of tasks
- Execution analysis
 - CPU utilisation, Worst-Case Execution Time,
 - problems
 - overrun, jitter
- Tasks with deadlines and priority-based scheduling
 - static and dynamic priority
- Tasks based on event-triggered FSM

TASK-BASED ARCHITECTURES

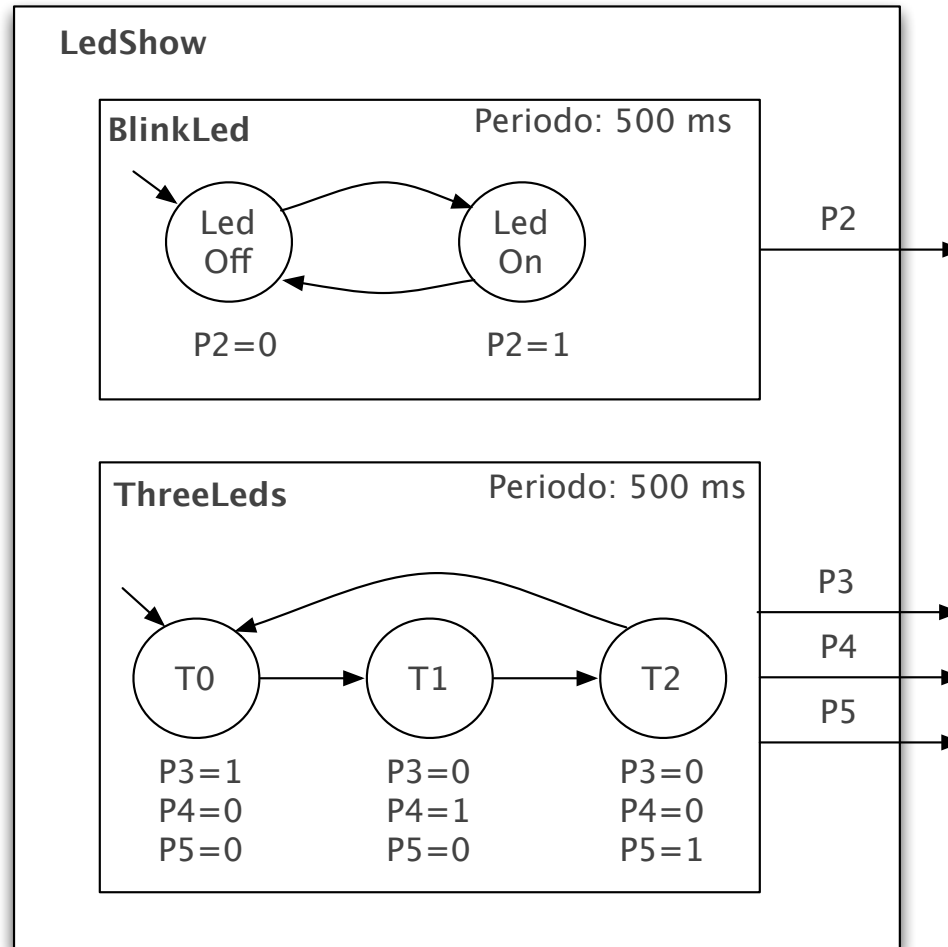
- Problem: complex embedded software modelling and design
 - need of proper approaches to decompose/modularise the behaviour and functionalities
- Approach: Task-based architectures
 - The behaviour of the embedded software is decomposed into a *set* of **concurrent tasks**
 - each task represents a specific well-defined confided unit of work/job to be done
 - the behaviour of *each* task can be described by a FSM
 - the global behaviour is the result of the execution and interaction of the concurrent FSMs

EXAMPLE: LED-SHOW

- LedShow example: 3 + 1 leds (p. 68, [PES])
 - One led: blinking with period: 500 ms
 - The other three leds: turn on/off in sequence, with a time interval of 500ms
- This could be modelled as a single task / FSM, however the modelling is much easier and simpler if we model it as the explicit composition of 2 tasks/FSM
 - single task/FSM version: much higher number of states and transitions

BLOCK DIAGRAM

- Representing each task as a block (rectangle), showing the input and output of each task/block:



TASK DECOMPOSITION: ADVANTAGES

- **Modularity**
 - every task is an independent module
 - the “interface” of the module in this case is the set of input/output variables / objects that the task is using (as input/output)
 - these can be shared with other modules
- Advantages
 - reduced complexity of the single task vs. the global behaviour
 - easier debugging
 - reusability

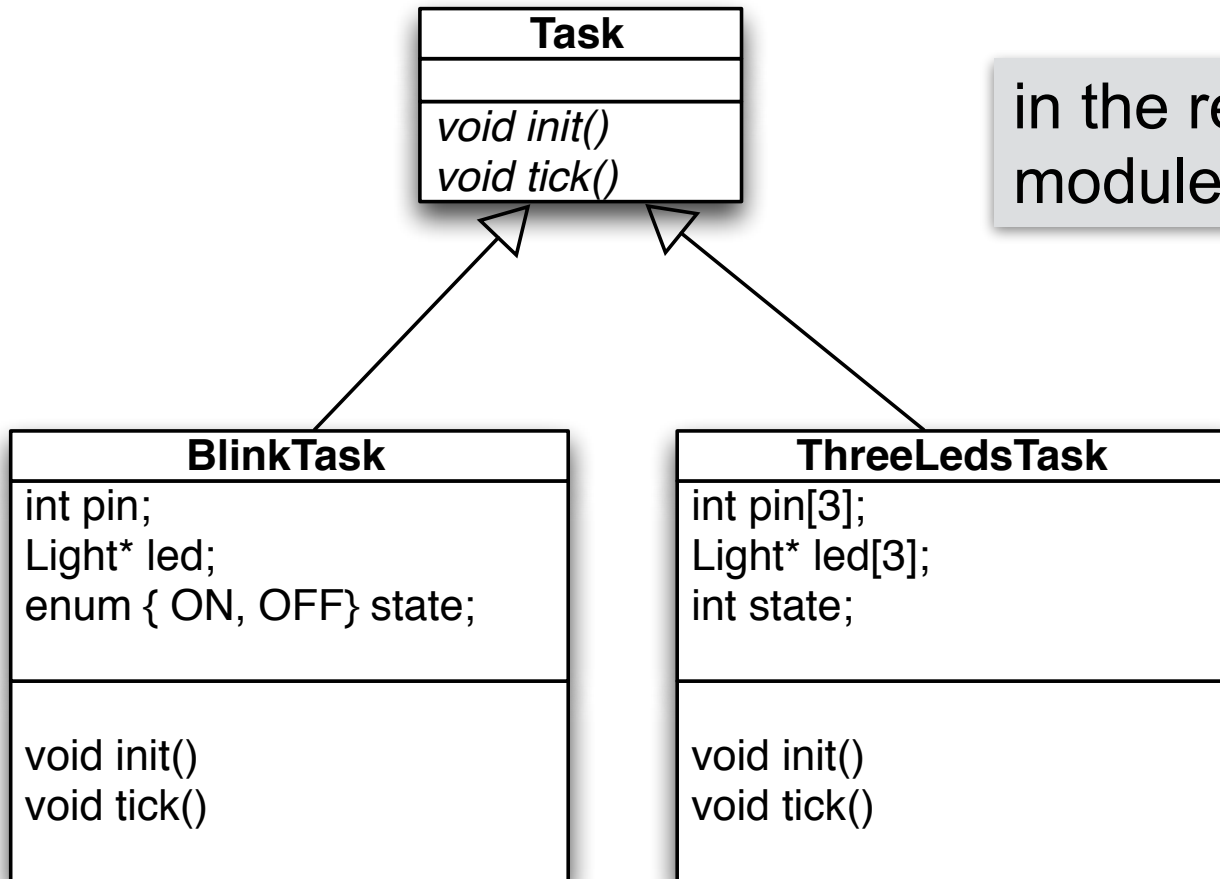
TASK DECOMPOSITION: CHALLENGES

- Tasks are **concurrent**
 - their execution overlaps in time
 - conceptually each task has its own logical control flow
- Tasks may have *dependencies* that create interactions among tasks and that need to be properly managed
 - interaction and coordination mechanisms
 - in our case, typically based on shared variables/objects

IMPLEMENTING TASK-BASED ARCHITECTURES - FIRST

- Introducing an abstract base class **Task**
 - **init** method
 - to initialise the task, called once
 - **tick** method (equivalent to the “step” method in FSM)
 - encapsulating the behaviour of the task
 - called periodically (period p)
- Each concrete tasks is implemented by extending this class
- Task execution carried on by periodically calling the tick method from the main loop

LED-SHOW EXAMPLE IN ARDUINO



in the repo:
[module-lab-2.3/led-show](https://github.com/arduino/module-lab-2.3/tree/main/led-show)

LED-SHOW EXAMPLE IN ARDUINO: TASK ABSTRACT CLASS

```
#ifndef __TASK__  
#define __TASK__  
  
class Task {  
  
public:  
    virtual void init() = 0;  
    virtual void tick() = 0;  
  
};  
  
#endif
```

LED-SHOW EXAMPLE IN ARDUINO: BLINK TASK

```
#ifndef __BLINKTASK__
#define __BLINKTASK__

#include "Task.h"
#include "Led.h"

class BlinkTask: public Task {

    int pin;
    Light* led;
    enum { ON, OFF} state;

public:

    BlinkTask(int pin);
    void init();
    void tick();
};

#endif
```

BlinkTask.h

```
#include "BlinkTask.h"

BlinkTask::BlinkTask(int pin){
    this->pin = pin;
}

void BlinkTask::init(){
    led = new Led(pin);
    state = OFF;
}

void BlinkTask::tick(){
    switch (state){
        case OFF:
            led->switchOn();
            state = ON;
            break;
        case ON:
            led->switchOff();
            state = OFF;
            break;
    }
}
```

BlinkTask.cpp

LED-SHOW EXAMPLE IN ARDUINO: THREE LEDS TASK

```
#ifndef __THREELEDSTASK__
#define __THREELEDSTASK__

#include "Task.h"
#include "Led.h"

class ThreeLedsTask: public Task {

    int pin[3];
    Light* led[3];
    int state;

public:

    ThreeLedsTask(int pin0, int pin1,
                  int pin2);

    void init();
    void tick();
};

#endif
```

ThreeLedsTask.h

```
#include "ThreeLedsTask.h"

ThreeLedsTask::ThreeLedsTask(int pin0, int pin1,
                             int pin2){

    this->pin[0] = pin0;
    this->pin[1] = pin1;
    this->pin[2] = pin2;
}

void ThreeLedsTask::init(){
    for (int i = 0; i < 3; i++){
        led[i] = new Led(pin[i]);
    }
    state = 0;
}

void ThreeLedsTask::tick(){
    led[state]->switchOff();
    state = (state + 1) % 3;
    led[state]->switchOn();
}
```

ThreeLedsTask.cpp

LED-SHOW EXAMPLE IN ARDUINO: MAIN LOOP

```
#include "Timer.h"
#include "BlinkTask.h"
#include "ThreeLedsTask.h"

Timer timer;

BlinkTask blinkTask(2);
ThreeLedsTask threeLedsTask(3,4,5);

void setup(){
    blinkTask.init();
    threeLedsTask.init();
    timer.setupPeriod(500);
}

void loop(){
    timer.waitForNextTick();
    blinkTask.tick();
    threeLedsTask.tick();
}

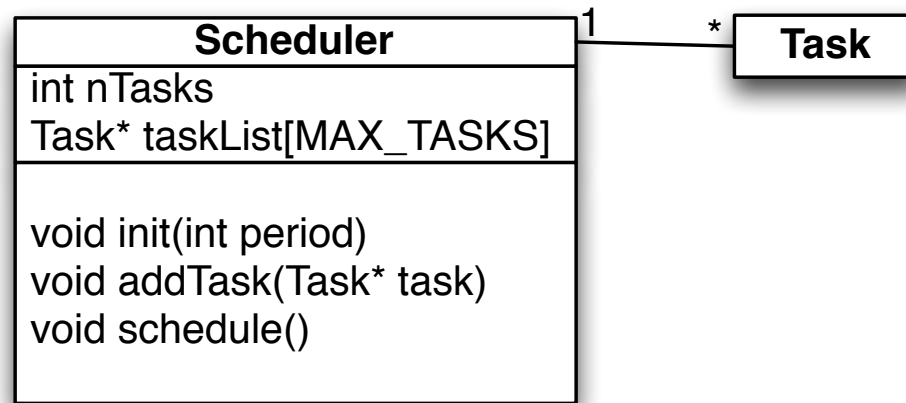
led-show.ino
```

MANAGING DIFFERENT PERIODS

- As soon as we want to manage tasks/FSMs with different periods, it is necessary:
 - to keep track of the specific period for each task
 - implement a task scheduling so that each task is called with its own period
- To this purpose, we implement a simple cooperative scheduler multi-tasking

SIMPLE COOPERATIVE SCHEDULER

- Scheduler
 - keeps track of a list of tasks to be executed



- Cooperative round-robin strategy
 - the scheduler has a period p equals to the *greatest common divisor* among the periods of the tasks
 - at each period p , it advances the execution of each tasks by calling its method `tick`
 - internally, it uses a timer to realise the periodic behaviour

A SIMPLE SCHEDULER

```
#ifndef __SCHEDULER__
#define __SCHEDULER__

#include "Timer.h"
#include "Task.h"

#define MAX_TASKS 10

class Scheduler {

    int basePeriod;
    int nTasks;
    Task* taskList[MAX_TASKS];
    Timer timer;

public:

    void init(int basePeriod);
    virtual bool addTask(Task* task);
    virtual void schedule();

};

#endif
```

Scheduler.h

```
#include "Scheduler.h"

void Scheduler::init(int basePeriod){
    this->basePeriod = basePeriod;
    timer.setupPeriod(basePeriod);
    nTasks = 0;
}

bool Scheduler::addTask(Task* task){
    if (nTasks < MAX_TASKS-1){
        taskList[nTasks] = task;
        nTasks++;
        return true;
    } else {
        return false;
    }
}

void Scheduler::schedule(){
    timer.waitForNextTick();
    for (int i = 0; i < nTasks; i++){
        if (taskList[i]->updateAndCheckTime(basePeriod)){
            taskList[i]->tick();
        }
    }
}
```


CLASS TASK REVISITED

```
#ifndef __TASK__
#define __TASK__

class Task {
    int myPeriod;
    int timeElapsed;

public:
    virtual void init(int period){
        myPeriod = period;
        timeElapsed = 0;
    }

    virtual void tick() = 0;

    bool updateAndCheckTime(int basePeriod){
        timeElapsed += basePeriod;
        if (timeElapsed >= myPeriod){
            timeElapsed = 0;
            return true;
        } else {
            return false;
        }
    }
};

#endif
```

LED-SHOW EXAMPLE

```
#include "Scheduler.h"
#include "BlinkTask.h"
#include "ThreeLedsTask.h"

Scheduler sched;

void setup(){

    sched.init(50);

    Task* t0 = new BlinkTask(2);
    t0->init(500);
    sched.addTask(t0);

    Task* t1 = new ThreeLedsTask(3,4,5);
    t1->init(150);
    sched.addTask(t1);

}

void loop(){
    sched.schedule();
}
```

LED-SHOW EXAMPLE

```
#ifndef __BLINKTASK__
#define __BLINKTASK__

#include "Task.h"
#include "Led.h"

class BlinkTask: public Task {

    int pin;
    Light* led;
    enum { ON, OFF} state;

public:

    BlinkTask(int pin);
    void init(int period);
    void tick();
};

#endif
```

BlinkTask.h

```
#include "BlinkTask.h"

BlinkTask::BlinkTask(int pin){
    this->pin = pin;
}

void BlinkTask::init(int period){
    Task::init(period);
    led = new Led(pin);
    state = OFF;
}

void BlinkTask::tick(){
    switch (state){
        case OFF:
            led->switchOn();
            state = ON;
            break;
        case ON:
            led->switchOff();
            state = OFF;
            break;
    }
}
```

BlinkTask.cpp

LED-SHOW EXAMPLE

```
#ifndef __THREELEDSTASK__
#define __THREELEDSTASK__

#include "Task.h"
#include "Led.h"

class ThreeLedsTask: public Task {

    int pin[3];
    Light* led[3];
    int state;

public:

    ThreeLedsTask(int pin0, int pin1,
                  int pin2);
    void init(int period);
    void tick();
};

#endif
```

ThreeLedsTask.h

```
#include "ThreeLedsTask.h"

ThreeLedsTask::ThreeLedsTask(int pin0, int pin1,
                             int pin2){

    this->pin[0] = pin0;
    this->pin[1] = pin1;
    this->pin[2] = pin2;
}

void ThreeLedsTask::init(int period){
    Task::init(period);
    for (int i = 0; i < 3; i++){
        led[i] = new Led(pin[i]);
    }
    state = 0;
}

void ThreeLedsTask::tick(){
    led[state]->switchOff();
    state = (state + 1) % 3;
    led[state]->switchOn();
}
```

ThreeLedsTask.cpp

REMARKS

- Implicit strong discipline to be adopted in task design
 - the execution of the tick method should have a duration always less than the scheduler period
- A different way to use timers and interrupts in schedulers concerns implementing the schedule step (schedule, tick) directly inside the interrupt/timer handler
 - **interrupt-driven schedulers**
 - main drawback
 - task execution would be inside interrupt handlers => constraints and limitations

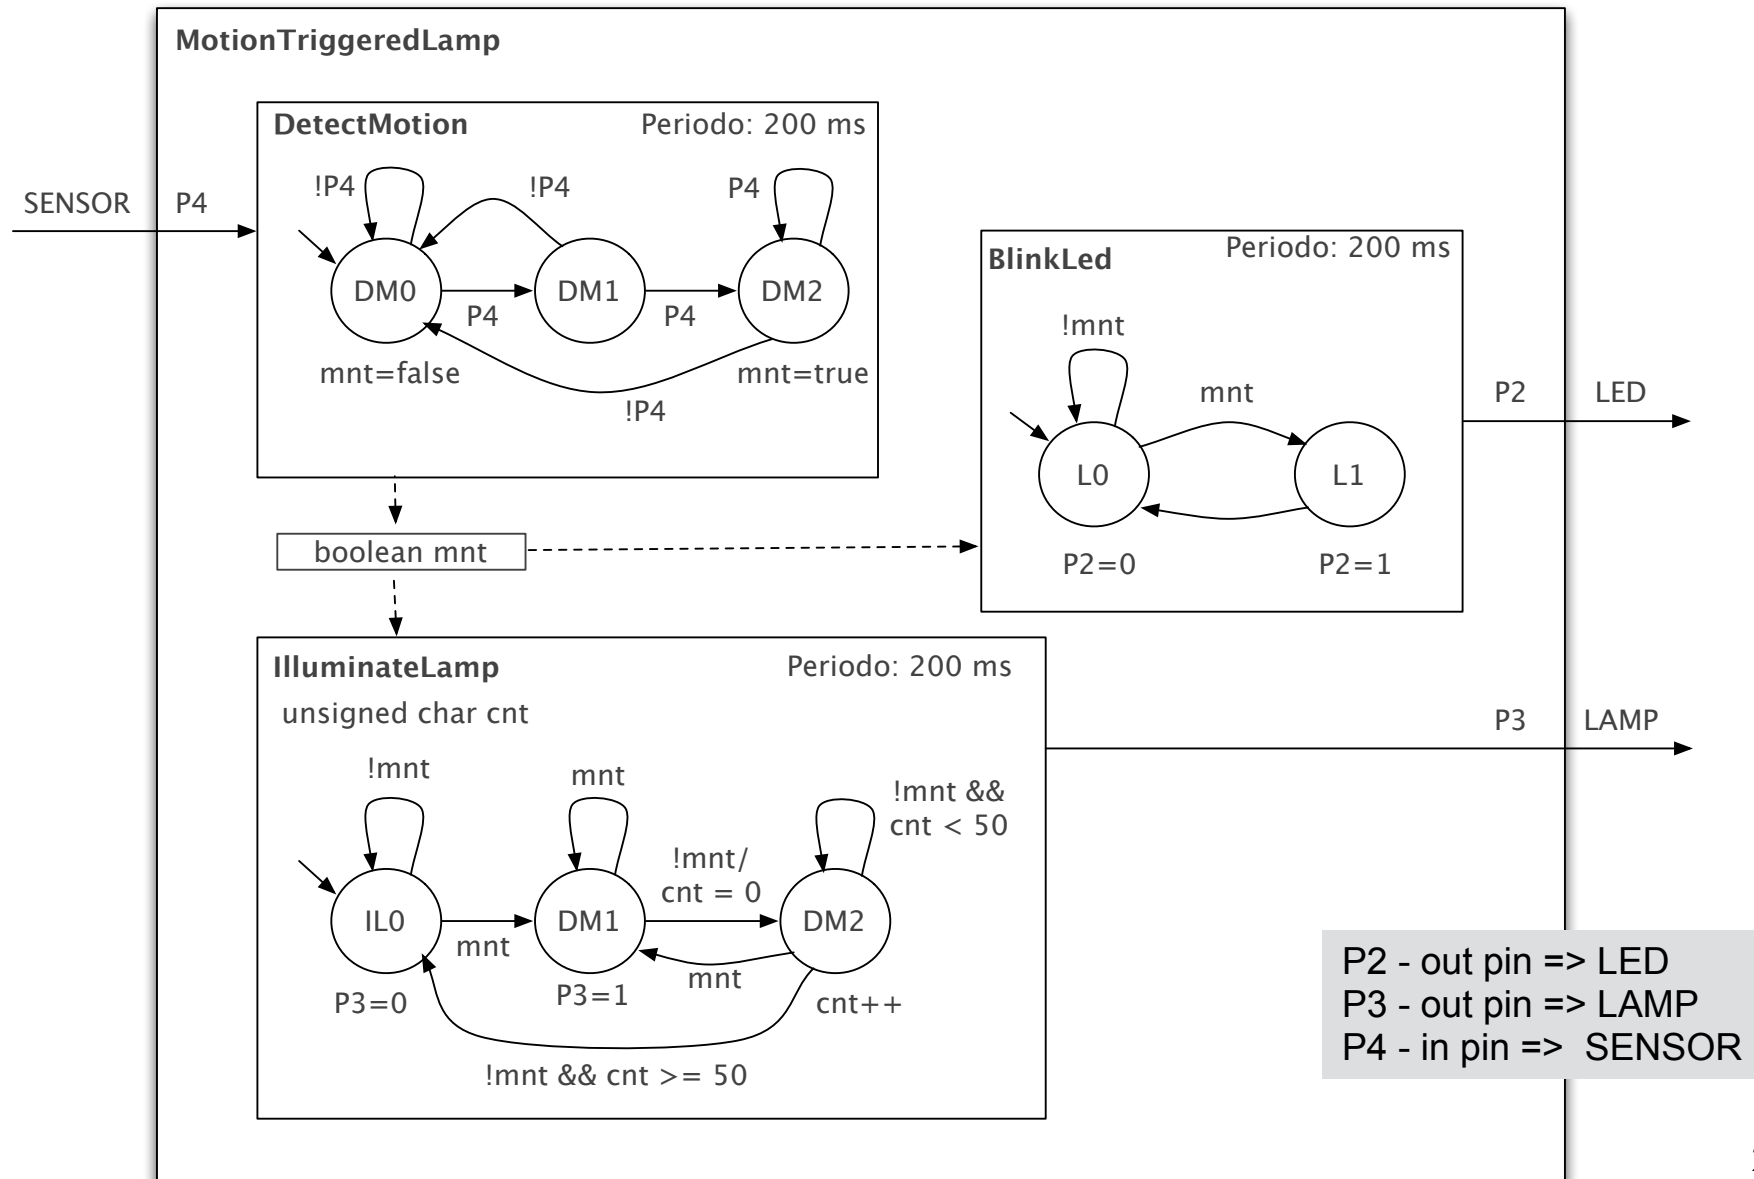
TASK DEPENDENCIES

- Task can have dependencies that requires various form of interactions
 - *temporal dependencies*
 - e.g.: a task T3 can be executed only after tasks T1 and T2
 - *producers/consumers*
 - e.g. a task T1 needs an info produced by a task T2
 - *data-oriented*
 - e.g. task T1 and T2 need to share some data
- In our case we represent and manage these dependencies through **shared variables**

EXAMPLE: MOTION-TRIGGERED LAMP SYSTEM

- Let's consider an embedded system in which a led is switched on as soon as some movement is detected, in some environment
 - it is equipped with a movement sensor, connected on the input pin P4
 - a movement is detected when it happens that P4 is sampled to be HIGH for two times in sequence, considering a period of 200ms.
 - when a movement is detected, a led connected to the output pin P4 should be switched on and keep the light on for 10 seconds, after the last detection time
 - the system includes a led, connected to P2, that should blink with period 200ms when a movement is detected (and keeps being detected)
- Design with three tasks
 - DetectMotion
 - IlluminateLamp
 - BlinkLed

EXAMPLE: MOTION-TRIGGERED LAMP SYSTEM



SHARED VARIABLES, ATOMIC ACTIONS AND RACE CONDITIONS

- Generally speaking, shared variables among concurrent tasks may lead to race conditions, in the case of concurrent reads/writes
- In our case (cooperative) there could not be races
 - each task tick is executed by the scheduler in sequence
 - the execution is atomic from the point of view of the whole system
- However the execution of multiple ticks of a task can be interleaved with the ones of other tasks
 - this may lead to high-level races, in some cases

EXPLOITING THE SLEEP MODE


- If the period of a scheduler is large enough, we can exploit the timer-driven behaviour and the sleep modality to save power
 - at each cycle, the scheduler after executing the task ticks, it goes to sleep until the next timer tick awakes it for the next cycle
- This is a very important extension for systems that are battery powered
 - nevertheless, in order to be applied, the scheduler period should be large enough to make irrelevant the latency of the microcontroller / MCU board to enter/exit the sleep modality

CPU USAGE AND SCHEDULING

- In synchronous FSM, we assume that actions have zero duration (i.e. instantaneous)
 - or: we abstract from their real duration
- Actually in the real system actions have always a duration and we need to carefully check that the chosen period is compatible with that duration, to avoid problems

OVERRUN EXCEPTION

- **Overflow exception** - when the execution time of actions exceeds the period
 - it is called a *timer overrun* if the scheduler is timer-based using interrupts (interrupt-driven schedulers)
 - in that case a new interrupt is generated before the conclusion of the interrupt handler
- Overflow example

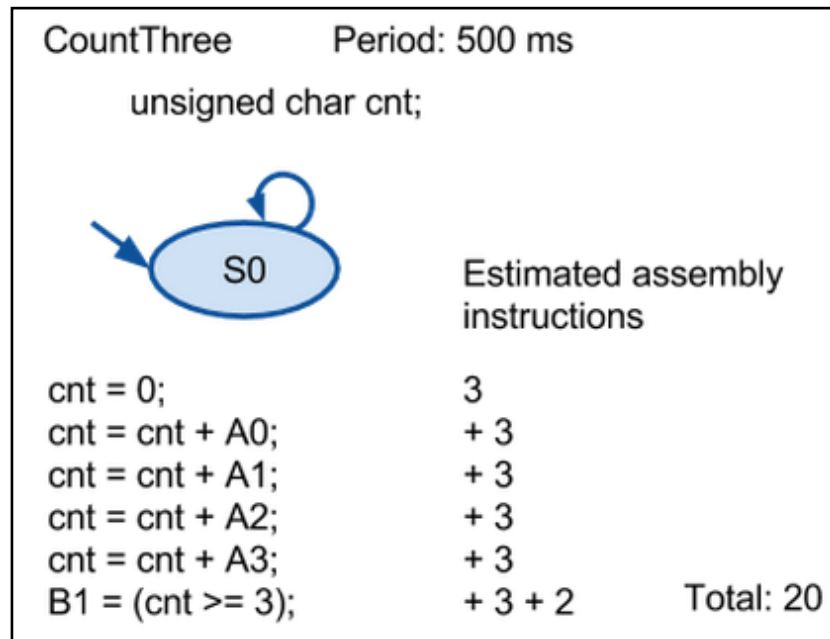


sync machine with period = 100ms

Depending on the state, different actions are executed. At $t = 200\text{ms}$, the duration of actions exceed the period...

CODE ANALYSIS

- Overrun exception can be spotted by doing the analysis of the code (assembly lines)
 - estimating the total duration of actions
 - check if in the worst case this duration exceeds the period
- Example ([PES], p. 101)
 - task CountThree with a single state



CPU UTILISATION PARAMETER AND WORST-CASE EXECUTION TIME

- The **CPU utilisation parameter** is the percentage of time in which the CPU (microcontroller) is used to execute a task:

$$U = (\text{time used for a task by the CPU} / \text{total time}) * 100\%$$

- The **Worst-Case-Execution-Time (WCET)** is the execution time in a period in the worst case
- In the case of multiple states/transitions, we consider the longest sequence of instructions

OVERRUN CASE


- We compute U . If U is $> 100\%$, then an overrun exception may occur
- In that case, in order to solve/avoid the problem, we can:
 - increase the period of the FSM
 - optimise the sequence of instructions to reduce the WCET
 - break long sequences in smaller sequences of actions
 - use a faster MCU
 - remove functionalities/behaviours from the system

CPU UTILISATION & WCET IN THE CASE OF MULTIPLE TASKS

- The analysis of CPU utilisation factor is even more frequent in systems with multiple tasks
- If tasks have the same period, then the WCET is computed by summing the individual WCETs of the tasks


LED-SHOW EXAMPLE

- LedShow example
 - BlinkLed (BL) and ThreeLeds (TL), each with a 500 ms period
 - Analysis example
 - suppose that the MCU would execute 100 instructions/sec, i.e. 0.01 sec/instruction
 - WCET for BL = 3 instructions $\Rightarrow 3 \cdot 0.01 = 0.03$ sec
 - WCET per TL = 9 instructions $\Rightarrow 9 \cdot 0.01 = 0.09$ sec
 - $U = (0.03 + 0.09) / 0.5$ (being 0.5 the period) = 24 %




LED-SHOW WITH OVERRUN

- Suppose that the WCET for BL is 200 ms and that the WCET of TL is 350
- Then: $U = (0.200+0.35)/0.5 = 1.1 \Rightarrow$ overrun



WHEN PERIODS ARE DIFFERENT

- In the general case, tasks can have *different periods*
- The WCET can be calculated by considering the **hyper-periods**:
 - periods that are the minimum common multiplier of the periods
- Example: LedShow - BL with period 300ms and TL period 200ms



- in 600 ms, BL executes $600/300 = 2$ times, mentre TL executes $600/200 = 3$ times
- the U parameter in the hyper-period is: $(2 \cdot 20\text{ms} + 3 \cdot 90\text{ms}) / 600\text{ms} = 55\%$

SUMMING UP

- Steps for the U calculation for $T_1 \dots T_n$ tasks over a microcontroller M
 - we determine the duration R of a single instruction of the MCU
 - we analyse the code of each T_i , computing its WCET
 - we first compute the max number of instructions in a tick and then we multiply it for R
 - we determine the hyperperiod H, as minimum common multiplier among the periods of all tasks ($T_1.H, T_2.H, \dots$)
 - U is then computed as:

$$U = ((H/T_1.\text{period}) * T_1.WCET + (H/T_2.\text{period}) * T_2.WCET + \dots) / H * 100$$

- $U > 100 \Rightarrow$ there will be overrun
- $U < 100 \Rightarrow$
 - for a single task: overrun cannot occur
 - multiple tasks: overrun can still occur

JITTER


- The **jitter** is the delay that occurs between the time in which a task is ready to be executed and the time in which the task is effectively executed
- Different scheduling strategies may lead to different jitters
 - LedShow example
 - if BL has priority => the jitter of BL is 0, while the jitter of TL is 30ms
 - if TL has priority => the jitter of BL is 90ms, while the jitter of TL is 0
- In general, giving priority to shortest tasks leads to minimise the average jitter
 - in the example: priority to BL

DEADLINE


- The **deadline** (*scadenza*) is defined as the interval of time within which a task *must* be executed after being ready to execute (in each period)
- If a task is not executed within its deadline, we have a ***missed-deadline exception*** that can result in a system failure
 - if a deadline is not specified, then it is its period, by default
- The scheduling strategy impacts on jitters and then also on the possibility to have missed-deadline exceptions

EXAMPLE

- LedShow example
 - suppose that BL has period = 300ms, WCET 30ms, and deadline 300ms
 - TL has period = 200ms, WCET 90ms and deadline 100ms
- If we give priority to BL, we have a missed deadline for TL:



- If we give priority to TL:



- In this case, we have a bigger jitter, but we don't have missed deadline

STATIC AND DYNAMIC PRIORITY SCHEDULING

- The **priority** concerns the order in which the tasks are executed
 - where there are multiple tasks that are ready, then the scheduler selects always the one with the higher priority
 - we can have both static and dynamic approaches in assigning the priority to tasks
 - static => priority assigned at design time and do not change at runtime
 - dynamic => priority assigned dynamically, according different strategies
- Key aspect in real-time OS (RTOS)
 - next modules

PERIODIC AND SPORADIC TASKS

- The tasks considered so far are **periodic**
 - they are executed periodically by the scheduler, given a period p , and this period corresponds to the period of the synchronous FSM
 - typically periodic tasks are static, that is: launched when the system is started and never removed
- Besides periodic tasks we can have **sporadic or aperiodic tasks**
 - i.e. tasks that are executed at arbitrary times

MANAGING APERIODIC TASKS

- Basic capabilities of a scheduler managing aperiodic tasks
 - dynamic insertion and removal of tasks
 - dynamic assignments of priorities
- Simple examples (in lab)
 - aperiodic task simulation with periodic tasks with “idle” state
 - simplest approach, not super efficient
 - tasks with active/not active meta-state
 - selected by a scheduler only if active
 - can be activated/de-activated by other tasks too

BIBLIOGRAPHY

- [PES] “Programming Embedded Systems - An introduction to Time-Oriented Programming”. Vahid et al.